

J-Sentinel Rule Engine Architecture

Executive Summary

The J-Sentinel rule engine, a critical component of the Analysis & Processing Layer, detects OWASP Top 10 vulnerabilities and custom logic flaws in Java applications by querying pre-stored code graphs and taint flows in Neo4j. It processes `codegraph.json` (code structure) and `taint_analysis.json` (tainted paths) using a Python-based rule engine with a YAML-based Domain-Specific Language (DSL), translated into Cypher queries.

1. Introduction

Purpose: The rule engine evaluates rules to identify vulnerabilities (e.g., log injection, SQL injection) in Java code, supporting J-Sentinel's secure DevSecOps-integrated static analysis. It queries pre-stored `codegraph.json` and `taint_analysis.json` data in Neo4j, enabling deep path analysis.

Objectives:

- Define the rule engine's architecture and components.
- Describe how it queries Neo4j-stored data.
- Plan for future CFG integration.
- Provide recommendations for implementation and testing.

Scope: The development team's responsibility is to build and test the rule engine, assuming `codegraph.json` and `taint_analysis.json` are already stored in Neo4j by another team.

2. Neo4j Data Model

The rule engine queries `codegraph.json` and `taint_analysis.json` data stored in Neo4j, structured as a graph with nodes and relationships.

2.1 Graph Schema

- **Nodes:**
 - `:FILE`: Source file (e.g., `test.java`; id: 1, name: 'test.java').
 - `:CLASS`: Class (e.g., `SimpleTest`; id: 2).

- `:METHOD`: Method (e.g., `run`; id: 3, name: 'run', parameters: 1).
- `:METHOD_CALL`: Method call (e.g., `logger.info`; id: 5, scope: 'logger', name: 'info').
- `:BINARY_EXPRESSION`: Expression (e.g., "Processing input: " + userInput; id: 11, operator: '+').
- `:SCAN`: Scan metadata (e.g., scanId: 'fd01a841-ff0e-4f1e-9c9c-8e01fe973ed8').
- Others: `:PARAMETER`, `:LOCAL_VARIABLE`, `:IF_STATEMENT`, `:STRING_LITERAL`, etc.
- **Relationships:**
 - `:CONTAINS`: Links containers (e.g., `:FILE` → `:CLASS`).
 - `:INVOKES`: Method calls (e.g., `:METHOD` → `:METHOD_CALL`).
 - `:DATA_FLOW`: Taint paths (e.g., `:METHOD_CALL {id: 6}` → `:METHOD_CALL {id: 5}`, with severity: 'HIGH').
 - `:PART_OF_SCAN`: Links nodes to `:SCAN` for scan-specific queries.
- **Properties:**
 - Nodes: id, name, type, scope, etc., from JSON.
 - `:DATA_FLOW`: severity, vulnerability, pathNodes (list of IDs).

2.2 Storage Format

- **codegraph.json:**
 - Nodes: Each `nodes` entry is a Neo4j node (e.g., `:METHOD_CALL {id: 5, name: 'info'}`).
 - Relationships: `edges` create relationships (e.g., `(:METHOD {id: 3})-[:INVOKES]->(:METHOD_CALL {id: 5})`).
 - Issues (e.g., `potentialLogInjections`) are node properties (e.g., `potentialLogInjection: 'Potential log injection'`).
- **taint_analysis.json:**
 - Nodes: Sources/sinks are `:METHOD_CALL` nodes (e.g., id: 6, name: 'readLine').
 - Relationships: `taintedPaths` create `:DATA_FLOW` (e.g., `(:METHOD_CALL {id: 6})-[:DATA_FLOW {severity: 'HIGH'}]->(:METHOD_CALL {id: 5})`).
 - Scan: Nodes link to `:SCAN` via `:PART_OF_SCAN`.

2.3 Querying Data

- **codegraph.json:**
 - Detects log injections (e.g., `methodCallId: 5`) by querying `:BINARY_EXPRESSION` with operator: '+' linked to `:METHOD_CALL {scope: 'logger'}`.

Example Cypher:

```
MATCH (method:METHOD)-[:CONTAINS_EXPRESSION]->(expr:BINARY_EXPRESSION
{operator: '+'})
```

```
MATCH (method)-[:INVOKES]->(call:METHOD_CALL {scope: 'logger', name: 'info'})
```

```
RETURN call, expr, call.potentialLogInjection
```

- **taint_analysis.json:**
 - Validates taint paths (e.g., `readLine` → `logger.info`; `id: 6` → `5`).

Example Cypher:

```
MATCH (source:METHOD_CALL {name:
'readLine'})-[:DATA_FLOW]->(sink:METHOD_CALL {scope: 'logger'})
```

```
RETURN source, sink, properties(r).severity
```

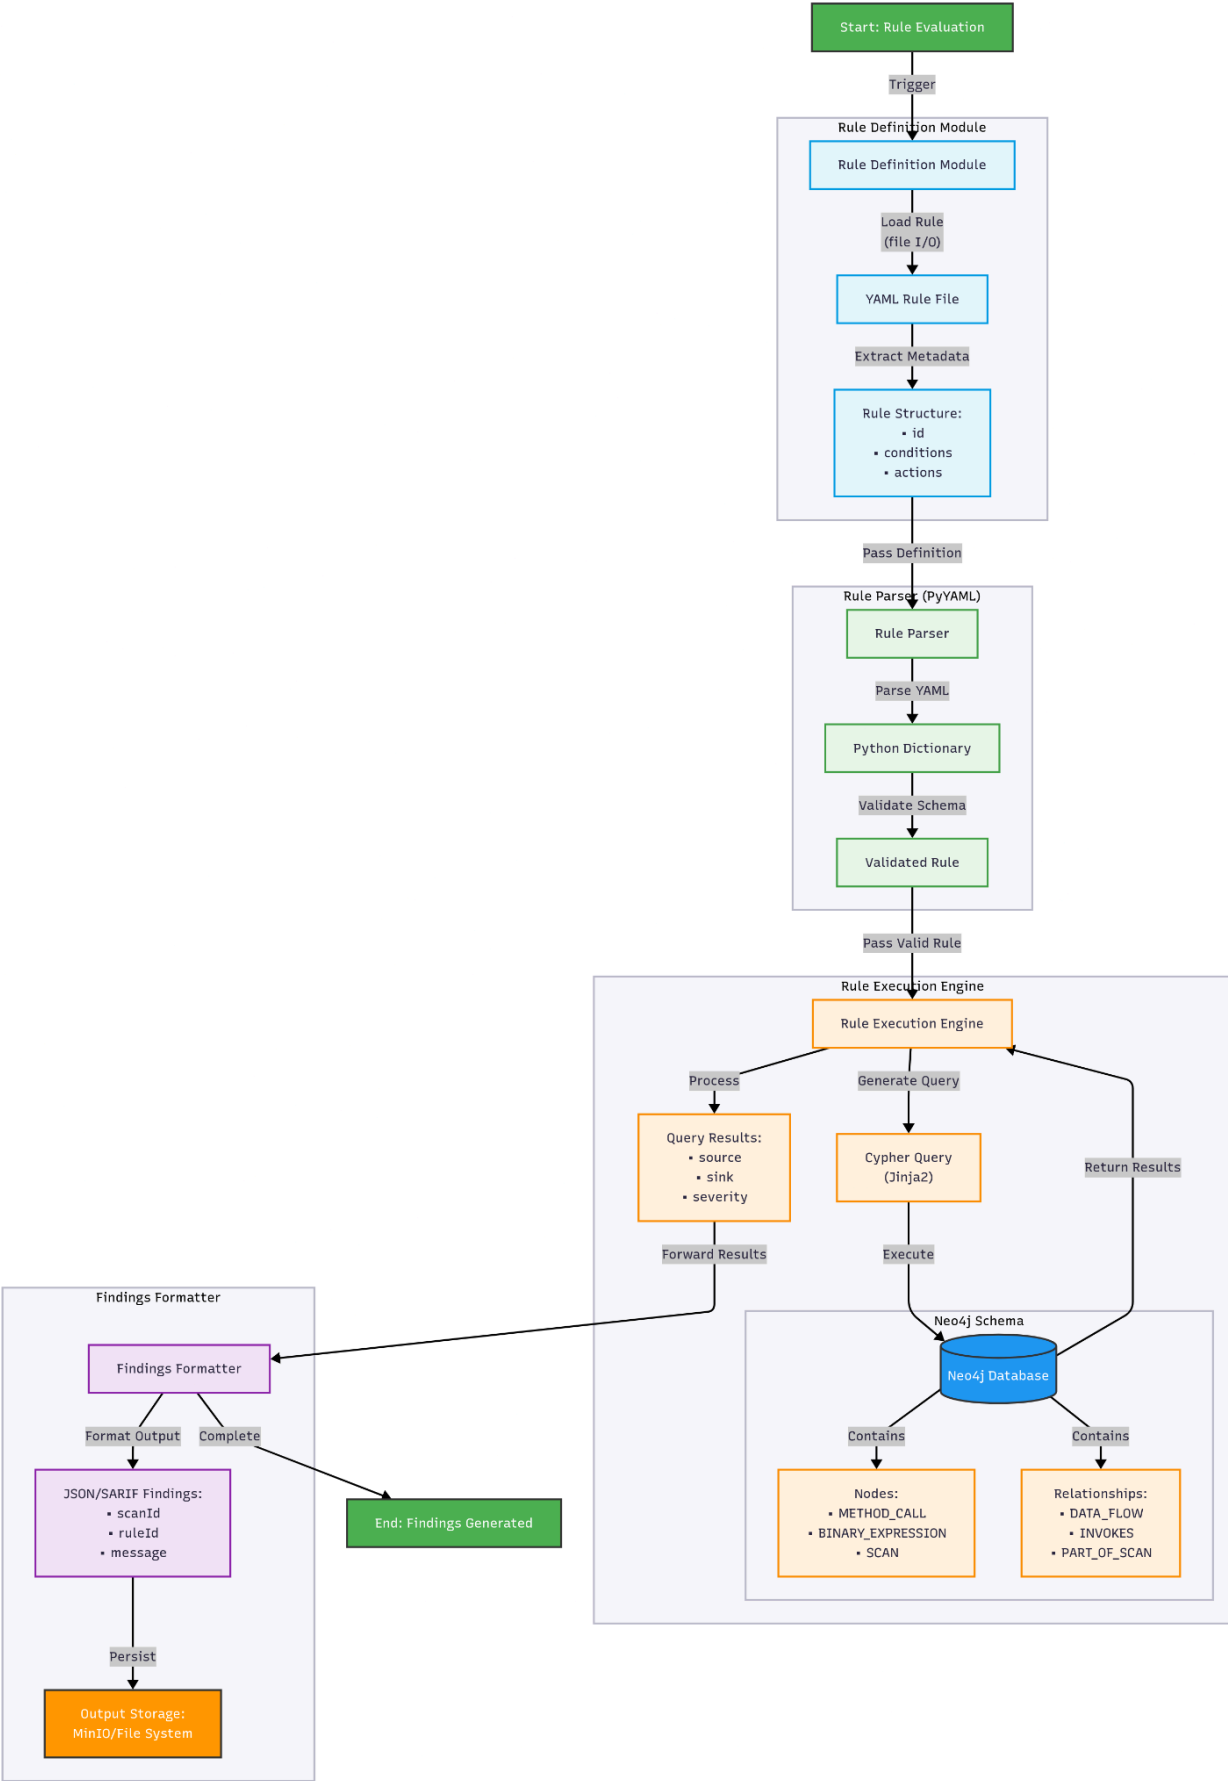
3. Rule Engine Architecture

The rule engine is a Python-based module that parses YAML rules, generates Cypher queries, and executes them against Neo4j. It integrates with J-Sentinel's Java-based CLI (`scanner.java`, `analyse.java`) via shared Neo4j data.

3.1 Tech Stack

- **Python 3.11:** Core language for rule parsing, query generation, and execution.
- **PyYAML 6.6:** Parses YAML rules.
- **Neo4j 5.0 with Cypher:** Queries stored data.
- **neo4j 5.16 (Python driver):** Executes Cypher queries.
- **Jinja2 3.1.4:** Generates Cypher queries.
- **pytest 8.3.3:** Tests rule engine components.

This diagram illustrates the end-to-end flow of the J-Sentinel rule engine, from YAML rule definition to vulnerability findings. It shows how rules are parsed, executed as Cypher queries against Neo4j, and formatted into JSON/SARIF reports. Arrows explicitly label each data transformation and handoff between components.



3.2 Components

3.2.1 Rule Definition Module

- **Function:** Defines rules using a YAML-based DSL.
- **Details:**
 - Rules specify vulnerability conditions (e.g., taint from `readLine` to `logger.info`).

Example rule (log injection):

rule:

id: LOG_INJECTION_001

name: Detect Log Injection

severity: HIGH

conditions:

- type: TAINT_PATH

source: { type: "METHOD_CALL", name: "readLine" }

sink: { type: "METHOD_CALL", scope: "logger", name: ["info", "severe"] }

actions:

- report:

message: "Potential log injection: {sink.name}"

suggested_fix: "Sanitize input with replaceAll('[\\n\\r]', '')"

- Stored in a file system or MinIO, accessed by the rule engine.

3.2.2 Rule Parser

- **Function:** Parses YAML rules into Python objects for query generation.
- **Details:**
 - Uses PyYAML to load YAML files.
 - Validates rule structure (e.g., presence of `id`, `conditions`).

Example:

```
import yaml
```

```
def parse_rule(yaml_file: str) -> dict:

    with open(yaml_file, 'r') as f:

        rule = yaml.safe_load(f)

    if not rule.get('id') or not rule.get('conditions'):

        raise ValueError("Invalid rule")

    return rule
```

3.2.3 Rule Execution Engine

- **Function:** Generates and executes Cypher queries against Neo4j.
- **Details:**
 - Uses Jinja2 to generate Cypher queries from rules.
 - Executes queries via neo4j Python driver.
 - Workflow:
 1. Parse YAML rule.
 2. Generate Cypher query (e.g., match `:DATA_FLOW` from `readLine` to `logger.info`).
 3. Query Neo4j for scan-specific data (via `:SCAN`).
 4. Collect findings (node IDs, paths, severity).

Example Cypher (generated):

```
MATCH path = (source:METHOD_CALL {name:
'readLine'})-[:DATA_FLOW*]->(sink:METHOD_CALL)

WHERE sink.scope = 'logger' AND sink.name IN ['info', 'severe']

RETURN source, sink, path, 'Potential log injection' AS message
```

3.2.4 Findings Formatter

- **Function:** Formats query results as JSON/SARIF.
- **Details:**
 - Outputs `scanId`, rule details, and findings.

Example:

```
{
```

```

"scanId": "fd01a841-ff0e-4f1e-9c9c-8e01fe973ed8",
"findings": [
  {
    "ruleId": "LOG_INJECTION_001",
    "severity": "HIGH",
    "message": "Potential log injection: info",
    "source": { "id": 6, "name": "readLine" },
    "sink": { "id": 5, "name": "info" },
    "suggestedFix": "Sanitize input with replaceAll('[\\n\\r]', '')"
  }
]
}

```

- Findings are saved to MinIO or a file system for downstream use.

4. Future CFG Integration

- **Control Flow Graph (CFG):**
 - **Nodes:** Statements (e.g., conditionals; `:CFG_NODE`).
 - **Relationships:** `:NEXT`, `:BRANCH`.
 - **Use:** Verify sanitization before sinks (e.g., check `replaceAll` before `logger.info`).
- **Implementation:**
 - Add `:CFG_NODE` labels to Neo4j.
 - Link to `:METHOD` via `:HAS_CFG`.

Update rules to traverse `:NEXT` paths:

```

MATCH (sanitize:CFG_NODE {name: 'replaceAll'})-[:NEXT*]->(sink:CFG_NODE {name:
'logger.info'})

```

```

RETURN sanitize, sink

```

5. Recommendations

1. Implement Core Components:

- Develop the rule parser using PyYAML.
- Build the query generator with Jinja2.
- Integrate neo4j driver for query execution.

2. Test Thoroughly:

- Write pytest tests for rule parsing, query generation, and execution.
- Validate against `taint_analysis.json` findings (e.g., log injection, `id: 6 → 5`).

Example test:

```
def test_parse_rule():
```

```
    rule = parse_rule("tests/log_injection.yaml")
```

```
    assert rule["id"] == "LOG_INJECTION_001"
```

○

3. Prepare for CFG:

- Design rules to query `:CFG_NODE` and `:NEXT` relationships.
- Test with sample CFG data when available.

4. Integration:

- Ensure the rule engine reads YAML rules from a shared directory or MinIO.
- Coordinate with the data loading team to verify Neo4j schema compatibility.

6. Conclusion

The J-Sentinel rule engine, implemented in Python, leverages Neo4j to detect vulnerabilities by querying pre-stored `codegraph.json` and `taint_analysis.json` data. Its modular architecture (parser, execution engine, formatter) and YAML-based DSL enable flexible rule definition and execution. Future CFG integration will enhance precision. The development team should focus on implementing the parser and query generator, testing against Neo4j data, and preparing for CFG support.